

PDF_Code

September 16, 2022

1 Imports

```
[ ]: import torch
import torch.nn as nn
from torchvision import datasets
import torchvision.transforms as tf
from torch.utils.data import DataLoader
import matplotlib.animation as animate
import matplotlib.pyplot as plt
import numpy as np
import torchvision.datasets
import random
```

2 Question 3

2.1 Create dataset/loader and create the random seed

```
[ ]: # seed = 393 # for same results
seed = random.randint(1, 10000) # for variation
print("Random Seed: ", seed)
random.seed(seed)
torch.manual_seed(seed)

img_size = 64
batch_size = 128
workers = 4
root="./data/CelebA"
nz = 100
# Size of feature maps in generator
gen_features = 64
# Size of feature maps in discriminator
disc_features = 64
# Number of training epochs
num_epochs = 1
# Learning rate for optimizers
D_lr = 0.002
G_lr = 0.001
```

```

dataset = datasets.ImageFolder(root=root,
                               transform=tf.Compose([
                                   tf.Resize(img_size),
                                   tf.CenterCrop(img_size),
                                   tf.ToTensor(),
                                   tf.Normalize(mean=(0.50612009, 0.42543493, 0.
↪38282761), std=(0.26589054, 0.24521921, 0.24127836)),
                               ]))

dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
↪shuffle=True, num_workers=workers)

DEVICE = 'mps' if torch.backends.mps.is_available() else 'cuda' if torch.cuda.
↪is_available() else 'cpu'

# Plot some training images
real_batch = next(iter(dataloader))
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")
grid_img = torchvision.utils.make_grid(real_batch[0])
plt.imshow(np.transpose(torchvision.utils.make_grid(real_batch[0][:64],
↪padding=5, normalize=True).cpu(),(1,2,0)))

```

Random Seed: 4120

[]: <matplotlib.image.AxesImage at 0x1482bdc10>

Training Images



2.2 Setup weight function

```
[ ]: def weights_init(m):  
    classname = m.__class__.__name__  
    if classname.find('Conv') != -1:  
        nn.init.normal_(m.weight.data, 0.0, 0.02)  
    elif classname.find('BatchNorm') != -1:  
        nn.init.normal_(m.weight.data, 1.0, 0.02)  
        nn.init.constant_(m.bias.data, 0)
```

2.3 Generator Setup and Initialization

```
[ ]: class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()
        self.gen = nn.Sequential(
            # N x nz x 1 x 1
            nn.ConvTranspose2d(in_channels=nz, out_channels=gen_features * 8,
            ↪kernel_size=4, stride=1, padding=0, bias=False),
            nn.BatchNorm2d(gen_features * 8),
            nn.Dropout(),
            nn.ReLU(True),
            # N x (gen_features*8) x 4 x 4
            # Paper uses kernel_size=5, stride=2 ==> cannot get 8x8 output
            nn.ConvTranspose2d(in_channels=gen_features * 8,
            ↪out_channels=gen_features * 4, kernel_size=4, stride=2, padding=1,
            ↪bias=False),
            nn.BatchNorm2d(gen_features * 4),
            # nn.Dropout(),
            nn.ReLU(True),
            # N x (gen_features*4) x 8 x 8
            nn.ConvTranspose2d(in_channels=gen_features * 4,
            ↪out_channels=gen_features * 2, kernel_size=4, stride=2, padding=1,
            ↪bias=False),
            nn.BatchNorm2d(gen_features * 2),
            # nn.Dropout(),
            nn.ReLU(True),
            # N x (gen_features*2) x 16 x 16
            nn.ConvTranspose2d(in_channels=gen_features * 2,
            ↪out_channels=gen_features, kernel_size=4, stride=2, padding=1, bias=False),
            nn.BatchNorm2d(gen_features),
            nn.ReLU(True),
            # N x gen_features x 32 x 32
            nn.ConvTranspose2d(in_channels=gen_features, out_channels=3,
            ↪kernel_size=4, stride=2, padding=1, bias=False),
            nn.Tanh()
            # N x 3 x 64 x 64
        )

    def forward(self, x):
        return self.gen(x)
```

```
[ ]: # Create the generator
netG = Generator().to(DEVICE)

netG.apply(weights_init)
```

```
print(netG)
```

```
Generator(
  (gen): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): ReLU(inplace=True)
    (4): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
    (5): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (6): ReLU(inplace=True)
    (7): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
    (8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (9): ReLU(inplace=True)
    (10): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
    (11): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (12): ReLU(inplace=True)
    (13): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1), bias=False)
    (14): Tanh()
  )
)
```

2.4 Discriminator setup and initialization

```
[ ]: class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.net = nn.Sequential(
            # Nx3x64x64
            nn.Conv2d(in_channels=3, out_channels=disc_features, kernel_size=4,
stride=2, padding=1, bias=False),
            nn.LeakyReLU(0.1),
            # N x disc_features x 32 x 32
            self._step(in_channels=disc_features, out_channels=disc_features *
2, kernel_size=4, stride=2, padding=1),
            # N x (disc_features*2) x 16 x 16
            self._step(in_channels=disc_features * 2,
out_channels=disc_features * 4, kernel_size=4, stride=2, padding=1),
```

```

        self._step(in_channels=disc_features * 4,
out_channels=disc_features * 8, kernel_size=4, stride=2, padding=1),
        # N x (disc_features*8) x 4 x 4
        nn.Conv2d(in_channels=disc_features * 8, out_channels=1,
kernel_size=4, stride=1, padding=0),
        nn.Sigmoid()
        # N x 1 x 1 x 1
    )

    def _step(self, in_channels, out_channels, kernel_size, stride, padding):
        return nn.Sequential(
            nn.Conv2d(in_channels=in_channels, out_channels=out_channels,
kernel_size=kernel_size, stride=stride, padding=padding, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.Dropout(),
            nn.LeakyReLU(0.1, inplace=True)
        )

    def forward(self, x):
        return self.net(x)

```

```

[ ]: # Create the Discriminator
netD = Discriminator().to(DEVICE)

netD.apply(weights_init)

print(netD)

```

```

Discriminator(
  (net): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
    (1): LeakyReLU(negative_slope=0.1)
    (2): Sequential(
      (0): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): Dropout(p=0.5, inplace=False)
      (3): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (3): Sequential(
      (0): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): Dropout(p=0.5, inplace=False)

```

```

        (3): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (4): Sequential(
      (0): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): Dropout(p=0.5, inplace=False)
      (3): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (5): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1))
    (6): Sigmoid()
  )
)

```

2.5 Initialize Loss and Optimizers

```

[ ]: criterion = nn.BCELoss()

# Create batch of latent vectors that we will use to visualize
# the progression of the generator
fixed_noise = torch.randn(64, nz, 1, 1, device=DEVICE)

real_label = 1.
fake_label = 0.

# Setup Adam optimizer for D and G
optimizerD = torch.optim.Adam(netD.parameters(), lr=D_lr, betas=(0.5, 0.999))
optimizerG = torch.optim.Adam(netG.parameters(), lr=G_lr, betas=(0.5, 0.999))

```

2.6 Training Loop

```

[ ]: # Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
for epoch in range(num_epochs):
    for i, data in enumerate(dataloader, 0):
        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        netD.zero_grad()
        real_cpu = data[0].to(DEVICE)
        b_size = real_cpu.size(0)

```

```

        label = torch.full((b_size,), real_label, dtype=torch.float,
↪device=DEVICE)
        output = netD(real_cpu).view(-1)
        errD_real = criterion(output, label)
        errD_real.backward()
        D_x = output.mean().item()

        noise = torch.randn(b_size, nz, 1, 1, device=DEVICE)
        fake = netG(noise)
        label.fill_(fake_label)
        output = netD(fake.detach()).view(-1)
        errD_fake = criterion(output, label)
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        lossD = errD_real + errD_fake
        optimizerD.step()

        #####
        # (2) Update G network: maximize log(D(G(z)))
        #####
        netG.zero_grad()
        label.fill_(real_label)
        output = netD(fake).view(-1)
        lossG = criterion(output, label)
        lossG.backward()
        D_G_z2 = output.mean().item()
        optimizerG.step()

        if i % 50 == 0:
            print(f"[{epoch+1}/{num_epochs}] [{i}/{len(dataloader)}] \tLoss_D:
↪{lossD.item():.4f} \tLoss_G: {lossG.item():.4f} \tD(x): {D_x:.4f} \tD(G(z)):
↪{D_G_z1:.4f}/{D_G_z2:.4f}")

            # save loss for plot
            G_losses.append(lossG.item())
            D_losses.append(lossD.item())

            # save G output on fixed noise
            if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i ==
↪len(dataloader)-1)):
                with torch.no_grad():
                    fake = netG(fixed_noise).detach().cpu()
                    img_list.append(torchvision.utils.make_grid(fake, padding=2,
↪normalize=True))

        iters += 1

```



Starting Training Loop...

[1/1] [0/1583]	Loss_D: 1.9389	Loss_G: 0.0054	D(x): 0.3618	D(G(z)):
0.3226/0.9947				
[1/1] [50/1583]	Loss_D: 2.0078	Loss_G: 3.9719	D(x): 0.5728	D(G(z)):
0.1172/0.0670				
[1/1] [100/1583]	Loss_D: 1.7966	Loss_G: 4.8704	D(x): 0.8013	D(G(z)):
0.6483/0.0214				
[1/1] [150/1583]	Loss_D: 1.0107	Loss_G: 2.2443	D(x): 0.6656	D(G(z)):
0.3240/0.1722				
[1/1] [200/1583]	Loss_D: 1.0700	Loss_G: 2.7194	D(x): 0.5891	D(G(z)):
0.2161/0.1431				
[1/1] [250/1583]	Loss_D: 1.5566	Loss_G: 1.0166	D(x): 0.3323	D(G(z)):
0.1381/0.4231				
[1/1] [300/1583]	Loss_D: 1.2354	Loss_G: 2.6069	D(x): 0.6985	D(G(z)):
0.4324/0.1294				
[1/1] [350/1583]	Loss_D: 1.1985	Loss_G: 2.2526	D(x): 0.8029	D(G(z)):
0.4995/0.1717				
[1/1] [400/1583]	Loss_D: 0.9133	Loss_G: 1.5177	D(x): 0.7235	D(G(z)):
0.2819/0.2728				
[1/1] [450/1583]	Loss_D: 1.5814	Loss_G: 1.7351	D(x): 0.3210	D(G(z)):
0.1007/0.2439				
[1/1] [500/1583]	Loss_D: 1.0627	Loss_G: 2.7073	D(x): 0.6374	D(G(z)):
0.3274/0.1183				
[1/1] [550/1583]	Loss_D: 1.0872	Loss_G: 1.5815	D(x): 0.5579	D(G(z)):
0.2525/0.2495				
[1/1] [600/1583]	Loss_D: 0.6390	Loss_G: 15.2075	D(x): 0.9767	
D(G(z)): 0.3870/0.0000				
[1/1] [650/1583]	Loss_D: 0.4340	Loss_G: 35.0539	D(x): 0.8366	
D(G(z)): 0.0000/0.0000				
[1/1] [700/1583]	Loss_D: 0.0128	Loss_G: 9.5067	D(x): 0.9923	D(G(z)):
0.0029/0.0042				
[1/1] [750/1583]	Loss_D: 0.0139	Loss_G: 12.1537	D(x): 0.9966	
D(G(z)): 0.0068/0.0043				
[1/1] [800/1583]	Loss_D: 0.0032	Loss_G: 11.5051	D(x): 0.9982	
D(G(z)): 0.0011/0.0008				
[1/1] [850/1583]	Loss_D: 0.0004	Loss_G: 10.8497	D(x): 0.9999	
D(G(z)): 0.0003/0.0003				
[1/1] [900/1583]	Loss_D: 0.0012	Loss_G: 10.0326	D(x): 0.9999	
D(G(z)): 0.0011/0.0020				
[1/1] [950/1583]	Loss_D: 0.0005	Loss_G: 11.6630	D(x): 0.9999	
D(G(z)): 0.0003/0.0010				
[1/1] [1000/1583]	Loss_D: 0.0010	Loss_G: 11.1038	D(x): 0.9994	
D(G(z)): 0.0004/0.0003				
[1/1] [1050/1583]	Loss_D: 0.0001	Loss_G: 12.1888	D(x): 0.9999	
D(G(z)): 0.0000/0.0001				
[1/1] [1100/1583]	Loss_D: 0.0003	Loss_G: 12.5416	D(x): 1.0000	
D(G(z)): 0.0002/0.0003				
[1/1] [1150/1583]	Loss_D: 0.0317	Loss_G: 13.4492	D(x): 0.9923	

D(G(z)): 0.0004/0.0000			
[1/1] [1200/1583]	Loss_D: 0.0026	Loss_G: 13.1606	D(x): 0.9983
D(G(z)): 0.0007/0.0002			
[1/1] [1250/1583]	Loss_D: 0.0002	Loss_G: 13.5758	D(x): 0.9999
D(G(z)): 0.0002/0.0002			
[1/1] [1300/1583]	Loss_D: 0.0006	Loss_G: 11.7854	D(x): 0.9997
D(G(z)): 0.0004/0.0005			
[1/1] [1350/1583]	Loss_D: 0.0003	Loss_G: 10.9696	D(x): 0.9998
D(G(z)): 0.0002/0.0002			
[1/1] [1400/1583]	Loss_D: 0.0005	Loss_G: 10.9147	D(x): 0.9997
D(G(z)): 0.0003/0.0001			
[1/1] [1450/1583]	Loss_D: 0.0004	Loss_G: 10.9153	D(x): 1.0000
D(G(z)): 0.0004/0.0007			
[1/1] [1500/1583]	Loss_D: 0.0003	Loss_G: 11.0962	D(x): 1.0000
D(G(z)): 0.0002/0.0001			
[1/1] [1550/1583]	Loss_D: 0.0001	Loss_G: 11.4345	D(x): 1.0000
D(G(z)): 0.0000/0.0003			

2.7 Visualise the loss

```
[ ]: plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses,label="G")
plt.plot(D_losses,label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



2.8 Animate the generators progression

```
[ ]: from IPython.display import HTML

fig = plt.figure(figsize=(8,8))
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0)), animated=True)] for i in img_list]
ani = animate.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000,
    ↪blit=True)

HTML(ani.to_jshtml())
```

```
[ ]: <IPython.core.display.HTML object>
```

